

CME 213, ME 339 Spring 2023

Introduction to parallel computing using MPI, openMP, and CUDA

Stanford University

Eric Darve, ICME, Stanford
Group Activity, May 22, 2023



Group activity

- These are instructions for the group activity.
- The goal of the activity is to calculate a prefix sum in parallel using a human computer.
- Let's review first how to calculate a reduction in parallel.

Reduction

Reduction is a classic problem in parallel computing.


A reduction is defined as the following computation:

$$s = \sum_{i=1}^n x_i$$

The goal is to compute this efficiently in parallel.

Binary tree

- The best strategy uses a binary tree.
- We start adding entries at the leaf and progressively go down the tree adding partial sums until we get the final result s .



Running time

- This approach is optimal in the sense that it minimizes the overall running time.
- For n entries, the total time is $O(\ln_2 n)$.

Prefix sum or scan

- The goal of the activity is to compute a prefix sum in parallel.
- Here is an example.
- Input: 3 5 6 2 4
- Output: 3 8 14 16 20


Mathematical definition of scan: $y_i = \sum_{j=1}^i x_j$

Algorithm 1

Algorithm 1 has 2 phases:

1. Collect or reduce partial results.
2. Distribute the partial sums in order to get the final result.

Algorithm 1




Complexity

- Number of passes: $2 \ln_2 n - 2$
- Flops is approximately 2x the amount of flops for a sequential algorithm.
- This algorithm is **work efficient**. It performs a small amount of extra work compared to the sequential algorithm.
- This extra work is required to make the algorithm parallel.
- This is a common tradeoff.

Algorithm 2

The Hillis and Steele algorithm performs more work by computing many reductions in parallel, but it can run faster in some cases.

Algorithm 2





**Many reductions are
computed in parallel**

Complexity

- Number of passes: $\ln_2 n$ vs. $2 \ln_2 n - 2$ for Algorithm 1.
- **Algorithm 2 is faster when many cores are available.**
- Flops is approximately $\ln_2 n$ the amount of flops for a sequential algorithm.
- **More work** is required compared to Algorithm 1.

Which algorithm is best?

- If you have fewer cores, Algorithm 1 is the best.
- If you have many cores, Algorithm 2 completes faster.
- In practice, hybrid algorithms can be designed that represent tradeoffs between reducing the number of stages in the algorithm (shorter span) and the number of flops (work efficiency).



Game time



Teams

- You should form teams of about 10 players.
- **Bring a lot of small 3in x 3in pieces of paper to use as scratch paper.**
- Goal of activity: parallel prefix sum using “human” processors

Human computers

- We are going to build a computer using humans.
- There are 3 types of players in this game: **mem, net, pu.**
- Each player has only one type.
- You can have as many players of each type as you want.

Rules!

- A **memloc** is a piece of paper with a single number written on it. You can strike a number to write a new one. A memloc is full when it has a number on it, and empty otherwise.
- There are 3 types of players: **mem**, **net**, **pu**. Each player has only one type. You can have as many players of each type as you want.
- **mem player:** can copy and strikethrough numbers. mem players are allowed to use scratch paper to organize the calculation.
- **net player:** can take/give memlocs to mem and pu. Cannot hold more than 3 memlocs at a time.
- **pu player:** takes two full memlocs and one empty memloc, adds the numbers on the full memlocs (may use a calculator) and writes the result on the empty memloc. Can take/give memlocs to net. Cannot hold more than 3 memlocs at a time.
- pu players must be at least 20 meters away from the mem players.

Generating the prefix scan problem

We will use the following code to generate sequence:

```
generate_sequence.cpp
```

See the notebook:

<https://colab.research.google.com/drive/1Zw8atH5szXmU6Uho--2fbWdzp7nwbGtJ?usp=sharing>

- Run the notebook.
- Enter your team number. The code returns a sequence of random numbers.
- This is the sequence you will use for the prefix scan.

Example: team 1 output

0s

1 !g++ generate_sequence.cpp

2 !group_number=1; echo \$group_number | ./a.out

Enter your group number (an integer greater or equal to 1)

Selected group ID: 1

Row 1: index

Row 2: random value

Index 1 to 10

1	2	3	4	5	6	7	8	9	10
57	65	28	84	10	33	43	78	86	27

Index 11 to 20

11	12	13	14	15	16	17	18	19	20
61	63	63	14	69	20	66	65	92	67

Index 21 to 30

21	22	23	24	25	26	27	28	29	30
41	72	53	50	83	37	17	20	93	83

Index 31 to 40

31	32	33	34	35	36	37	38	39	40
41	83	67	13	62	35	89	70	57	18

READY

SET

GO

Follow-up discussion

- How did you organize your group?
- What was the best strategy?
- What were the main bottlenecks?
- What would you do differently?